CSCE-689, Programming Assignment #2 Sentiment Analyzer
Due: Sunday, Feb. 19, 2017 by 11:00pm

Your goal for this homework is to perform Sentiment Analysis: classifying an entire movie review as positive or negative.

Train a Naive Bayes classifier and a Perceptron classifier on the imdb1 data set provided with the starter package. This is the actual Internet Movie Database review data used in the original Pang and Lee paper:

Bo Pang, Lillian Lee, and Shivakumar Vaithyanathan. 2002. Thumbs up? Sentiment Classification using Machine Learning Techniques. Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP), pp. 79-86.

The python starter code comes already set up for 10-fold cross-validation training and testing on this data. Recall that cross-validation involves dividing the data into several sections (10 in this case), then training and testing the classifier repeatedly, with a different section as the held-out test set each time. Your final accuracy is the average of the 10 runs. When using a movie review for training, you use the fact that it is positive or negative (the hand-labeled "true class") to help compute the correct statistics. But when the same review is used for testing, you only use this label to compute your accuracy. The data comes with the crossvalidation sections; they are defined in the file:

```
data/poldata.README.2.0
```

Naive Bayes Algorithm:

(3 Points!) Task 1- Your first task is to implement the Naive Bayes classifier training and testing code and evaluate them using the cross-validation mechanism.

(1 Point!) Task 2- Next, evaluate your model again with the stop words removed. Does this approach affect average accuracy (for the current given data set)?

(2 Points!) Task 3- Next, you will implement a binarized version of the multinomial Naive Bayes classifier that uses feature presence/absence (binary) instead of feature counts.

Two option booleans FILTER_STOP_WORDS, BOOLEAN_NB have been defined. If FILTER_STOP_WORDS is true, stop word tokens will be removed from both training and test documents. If the BOOLEAN_NB flag is true, then your code (specifically the train and evaluate functions) should train and test on the binarized version of the Naive Bayes classifier (with boolean features - for presence/absence of each word as opposed to word counts).

Perceptron Algorithm:

(4 Points!)

Next, implement the Perceptron classification algorithm. Your perceptron code should implement parameter averaging. Instead of using parameters by the end of training to classify test examples, use the average of the parameters after updating with each data instance in each iteration. A nice trick for doing this efficiently is described in section 2.1.1 of Hal Daumes thesis (http://www.umiacs.umd.edu/ hal/docs/daume06thesis.pdf).

The only hyperparameter is the number of iterations. Run the classifier and report the cross validation results with various numbers of iterations (for exaple: 1, 10, 50, 100, 1000).

1. Python Starter Code

You are encouraged to use Python for the programming assignments in this class, but if you prefer, you can use Java too. However, simple starter code will be provided in some of the programming assignments (including this one) and only in Python.

In the same directory where the provided code and data are, you can run the code:

```
cd python
python NaiveBayes.py ../data/imdb1
OR
python NaiveBayes.py -f ../data/imdb1
OR
python NaiveBayes.py -b ../data/imdb1
OR
python Perceptron.py ../data/imdb1/ 1 (run for 1 iteration)
```

The functions

```
def classify(self, words):
AND
def addExample(self, klass, words):
AND
def train(self, split, iterations):  (only in the code for Perceptron)
```

have been flagged for you with the universal "TODO" marker. You're free to add other elements further invoked from these three functions (you would need to add some data structures for your work - where to add them will be mentioned in the starter code).

OUTPUT FORMATTING

The results will look something like the figure in the next page. Even if you program using a language different from Python and do not use the starter code, your program output should follow the same structure.

GRADING CRITERIA

Your program will be graded all based on cross validation results on the provided training set, and mainly based on the correctness of your code (please refer back to the first section for how we will weight each task). We understand that you may produce slightly different results if you have implemented some details in a different way. Therefore, please describe enough details in your written report. In addition, we will test your code on a separate test set to make sure it runs properly. For instance, if we found any experimental result has been hard coded, you won't get any credit for that part.

ELECTRONIC SUBMISSION INSTRUCTIONS
(a.k.a. "What to turn in and how to tun in")

You need to submit 2 things:

1. The source code files for your program. Be sure to include <u>all</u> files that we will need to compile and run your program!

2. A report file that includes the following information:

   - how to compile and run your code
   - results and analysis
   - any known bugs, problems, or limitations of your program

   <u>REMINDER:</u> your program **must** compile and run. We will not grade programs that cannot be run.

How to turn in:

1. please include everything in a single package, have it compressed and emailed to our grader Guanlun at "guanlun@tamu.edu".

2. please wrtie the subject line as: "PA #2 Sentiment Analyzer for CSCE689 - UIN - LastName", please replace the placeholders "UIN" and "LastName" with your own UIN and LastName.

[INFO] Fold 0 Accuracy: 0.500000

[INFO] Fold 1 Accuracy: 0.500000

[INFO] Fold 2 Accuracy: 0.500000

[INFO] Fold 3 Accuracy: 0.500000

[INFO] Fold 4 Accuracy: 0.500000

[INFO] Fold 5 Accuracy: 0.500000

[INFO] Fold 6 Accuracy: 0.500000

[INFO] Fold 7 Accuracy: 0.500000

[INFO] Fold 8 Accuracy: 0.500000

[INFO] Fold 9 Accuracy: 0.500000

[INFO] Accuracy: 0.500000

Figure 1: Sample Output for Sentiment Analyzer